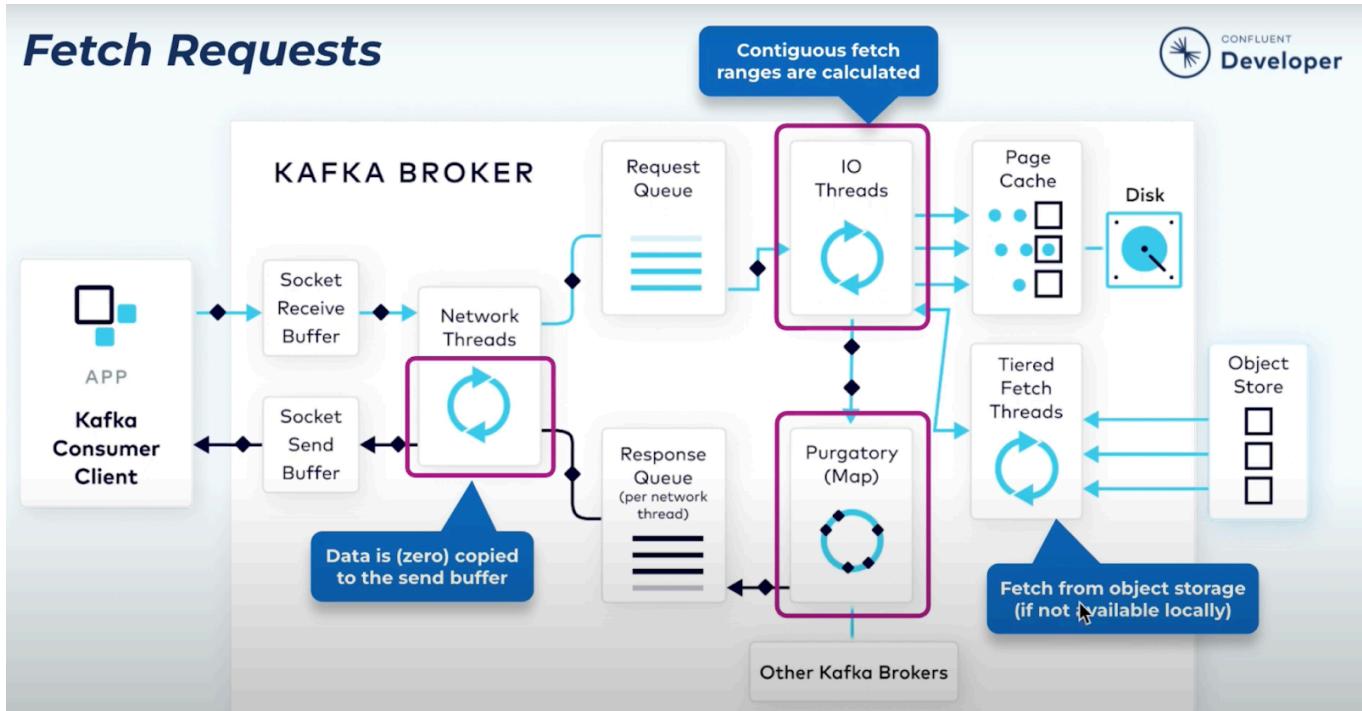
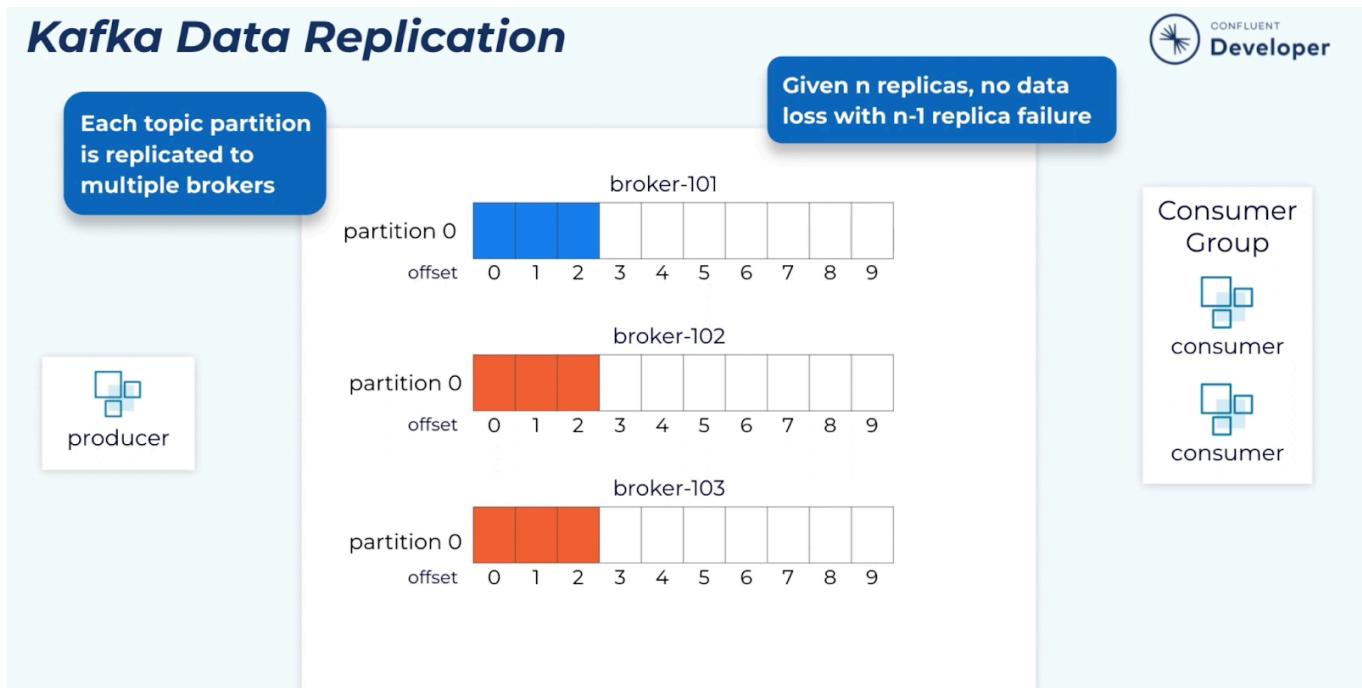


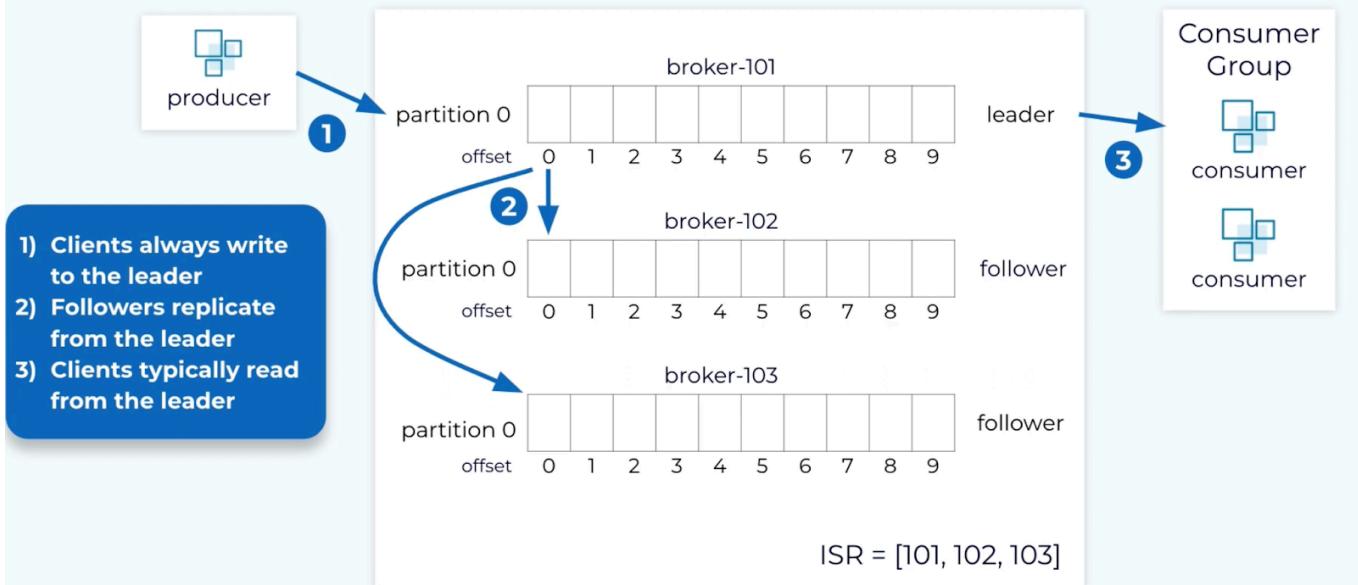
Kafka Architecture



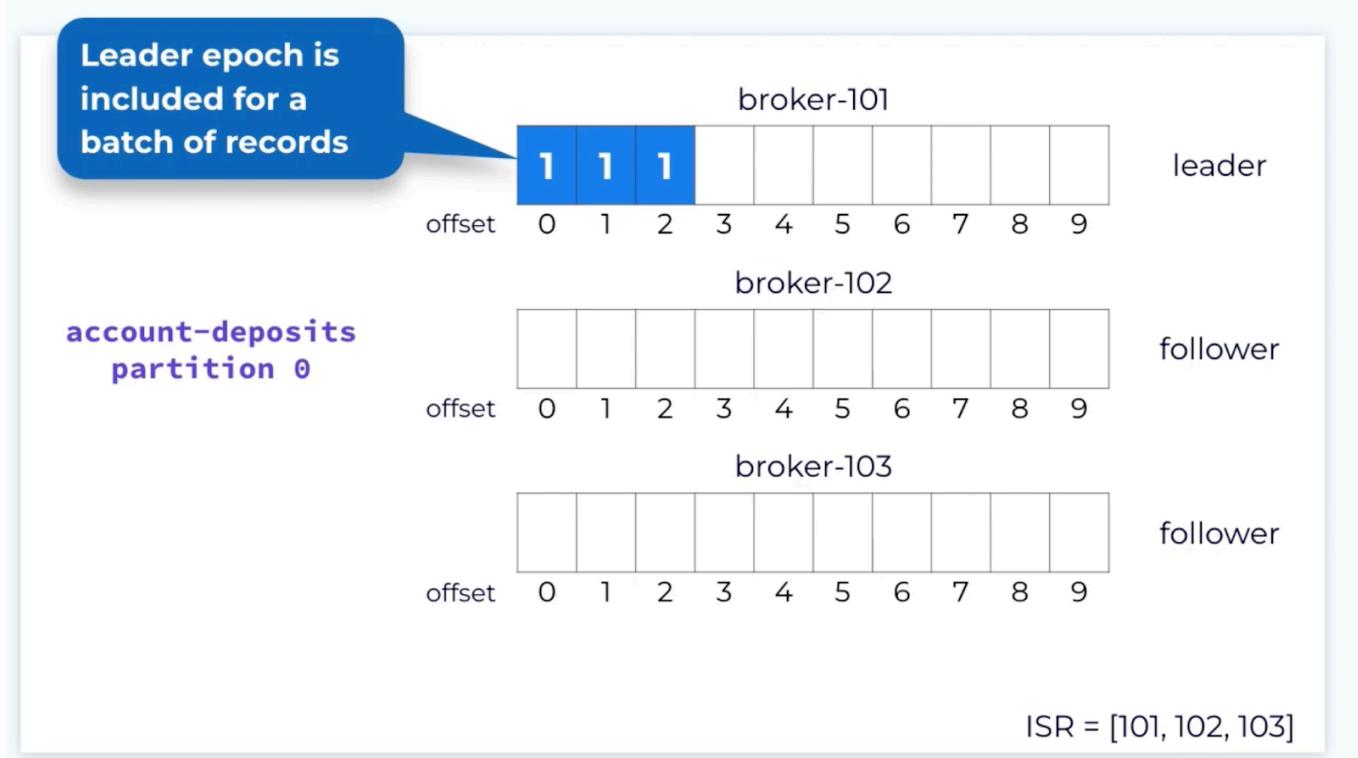
Kafka Data Replication:



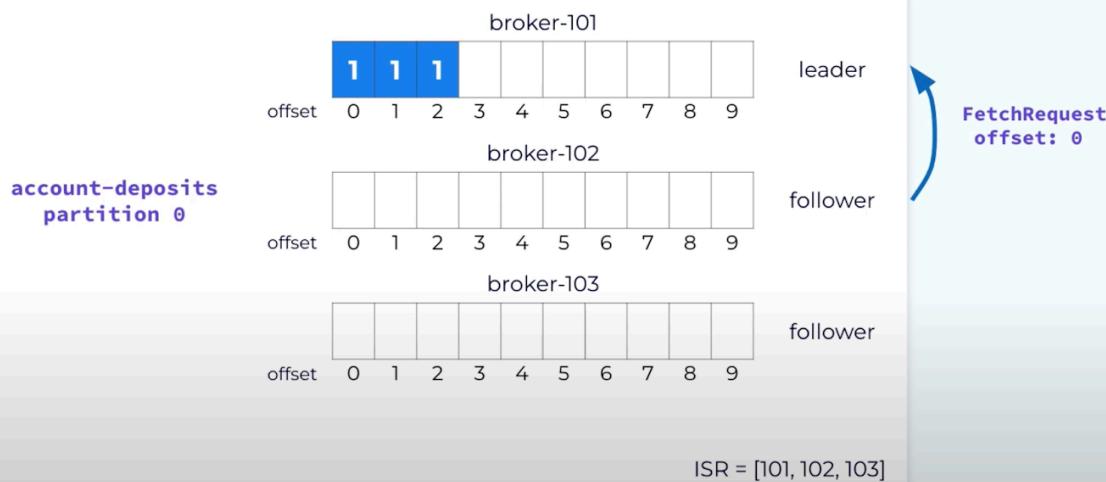
Leader, Follower, and In-Sync Replica (ISR) List



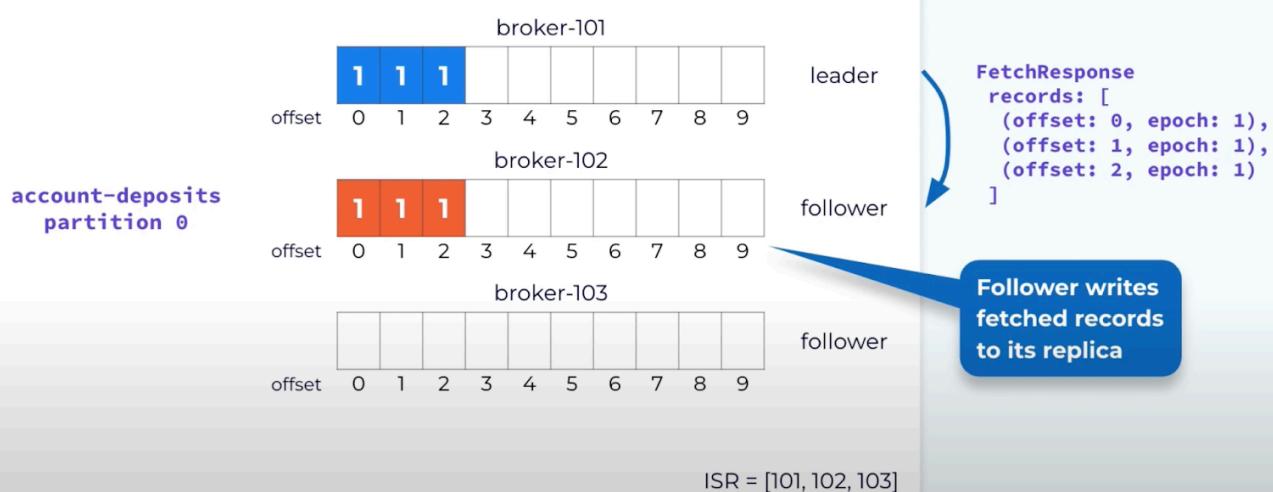
Leader Epoch



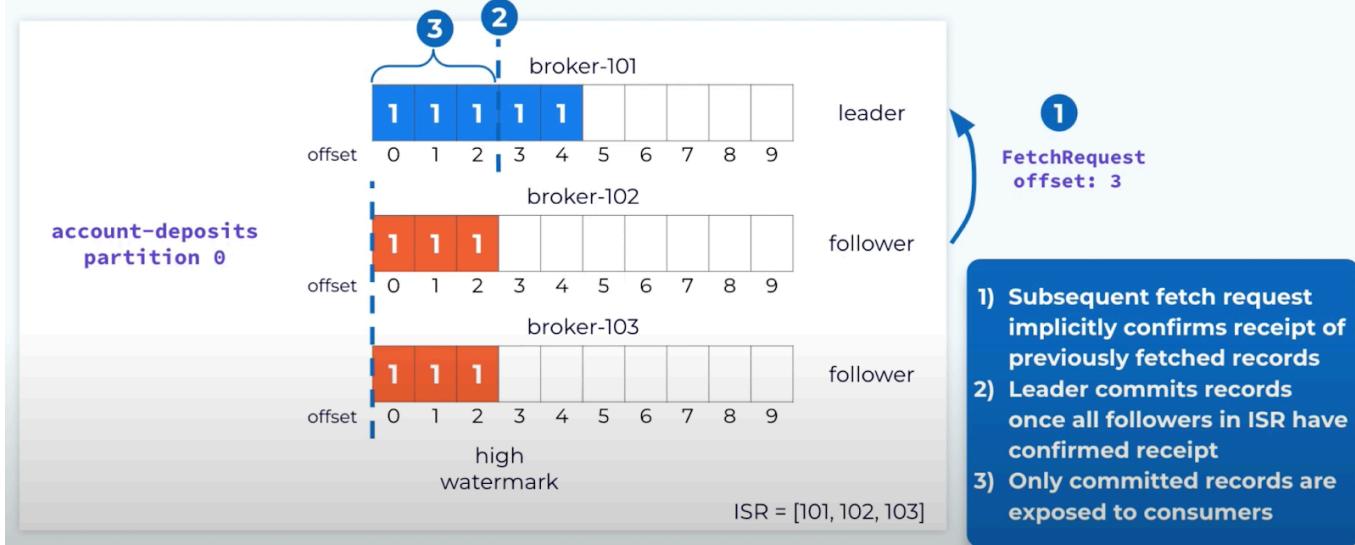
Follower Fetch Request



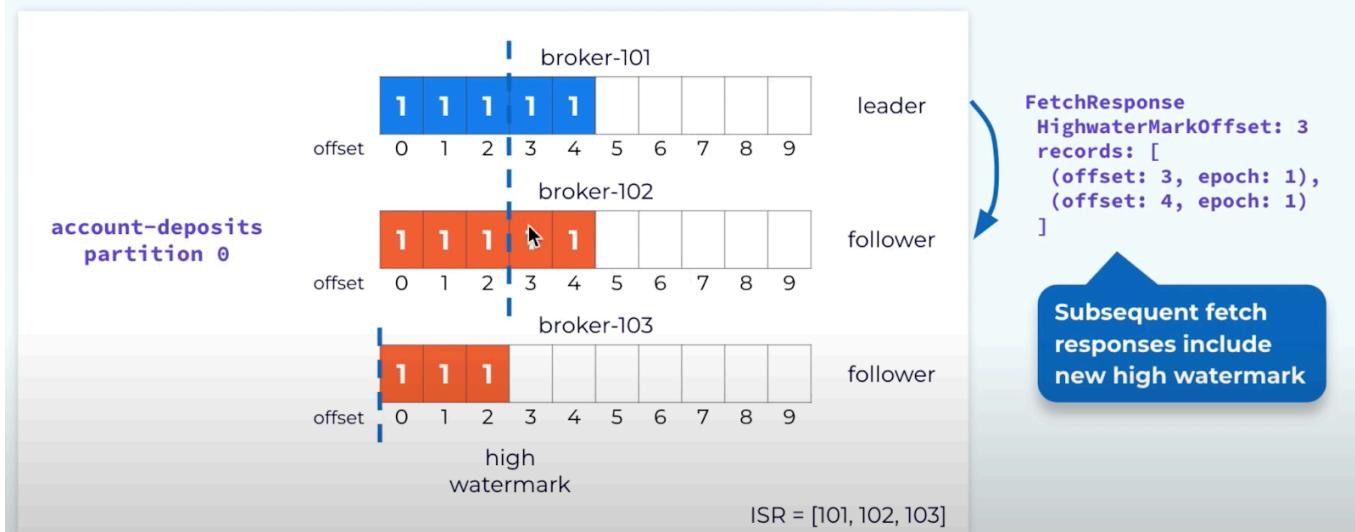
Follower Fetch Response



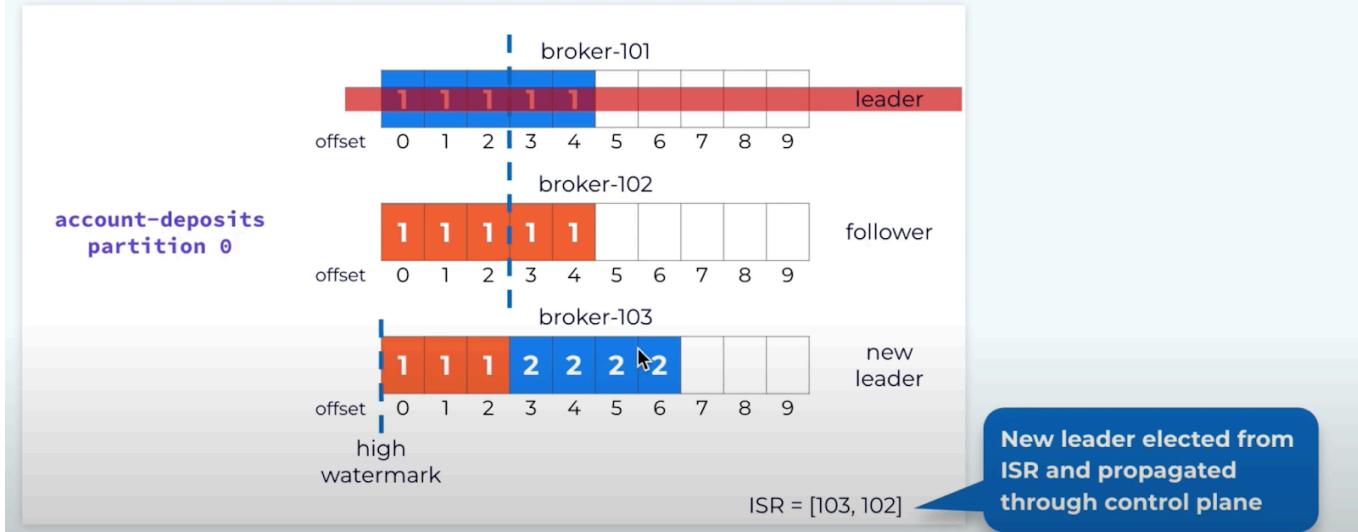
Committing Partition Offsets



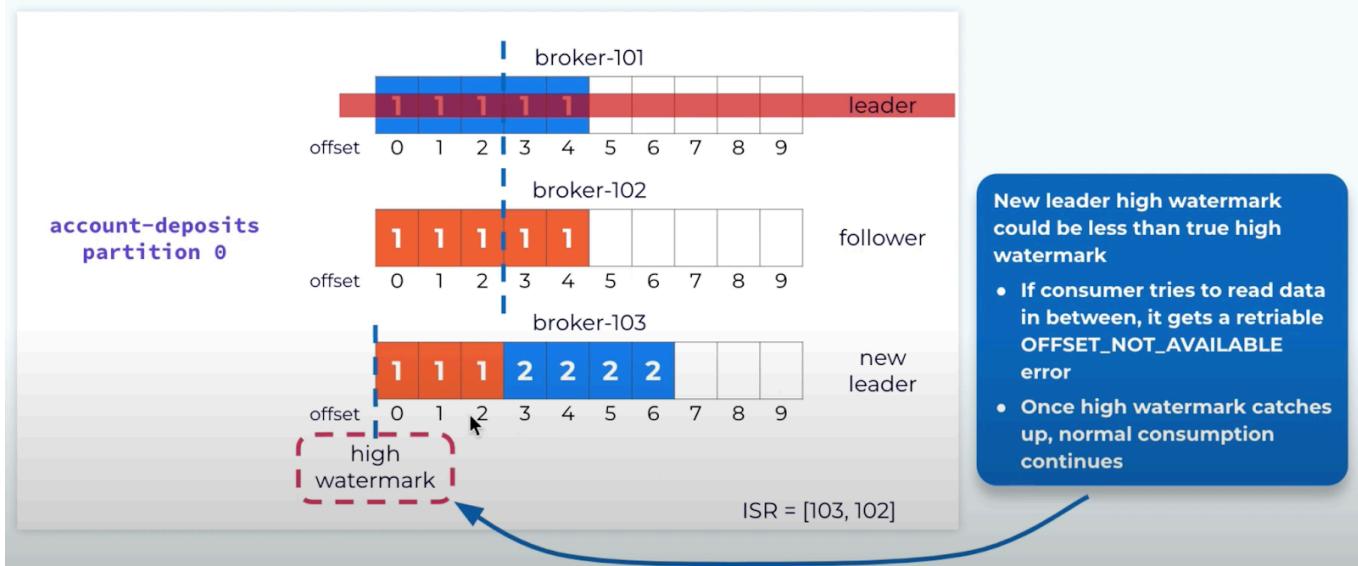
Advancing the Follower High Watermark



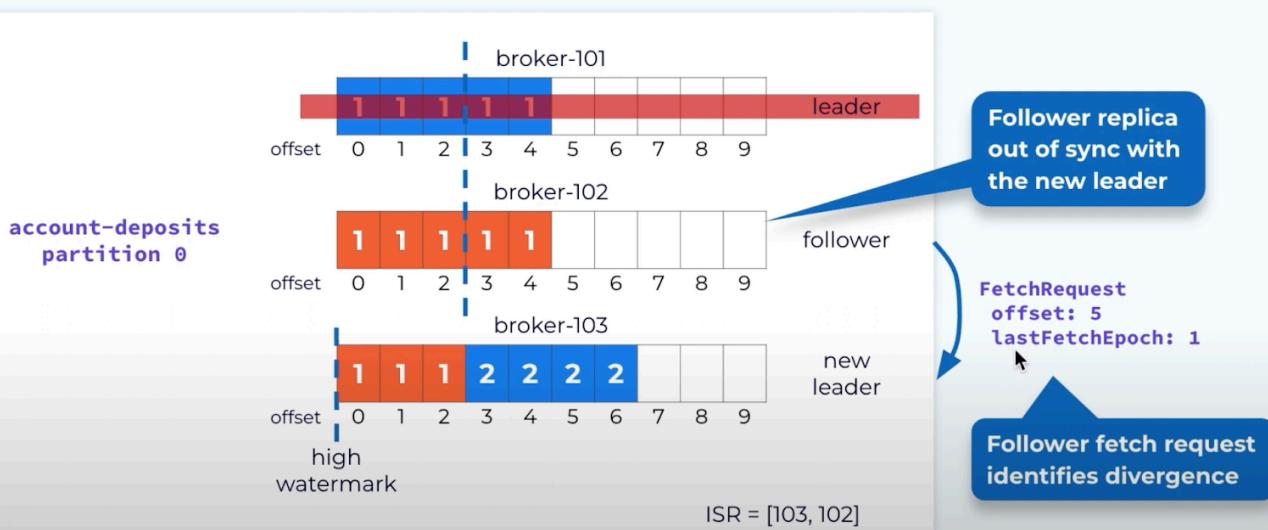
Handling Leader Failure



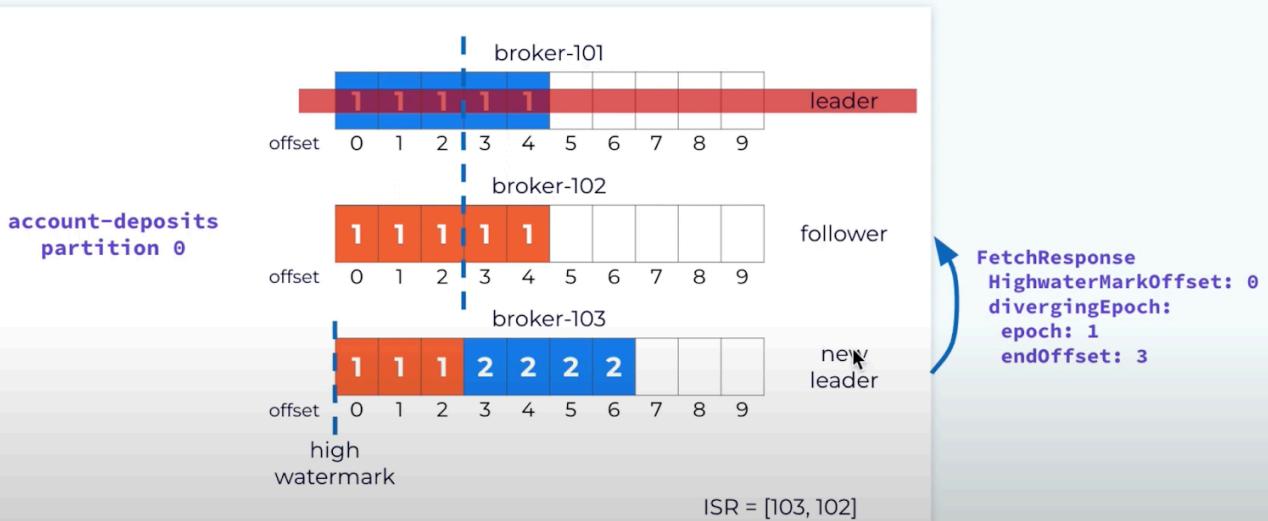
Temporary Decreased High Watermark



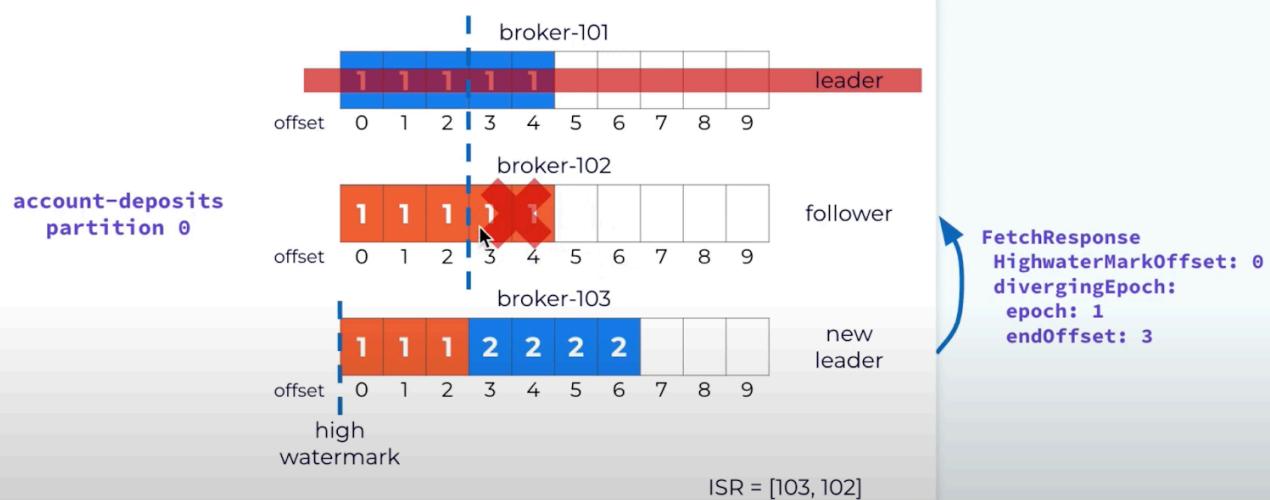
Partition Replica Reconciliation



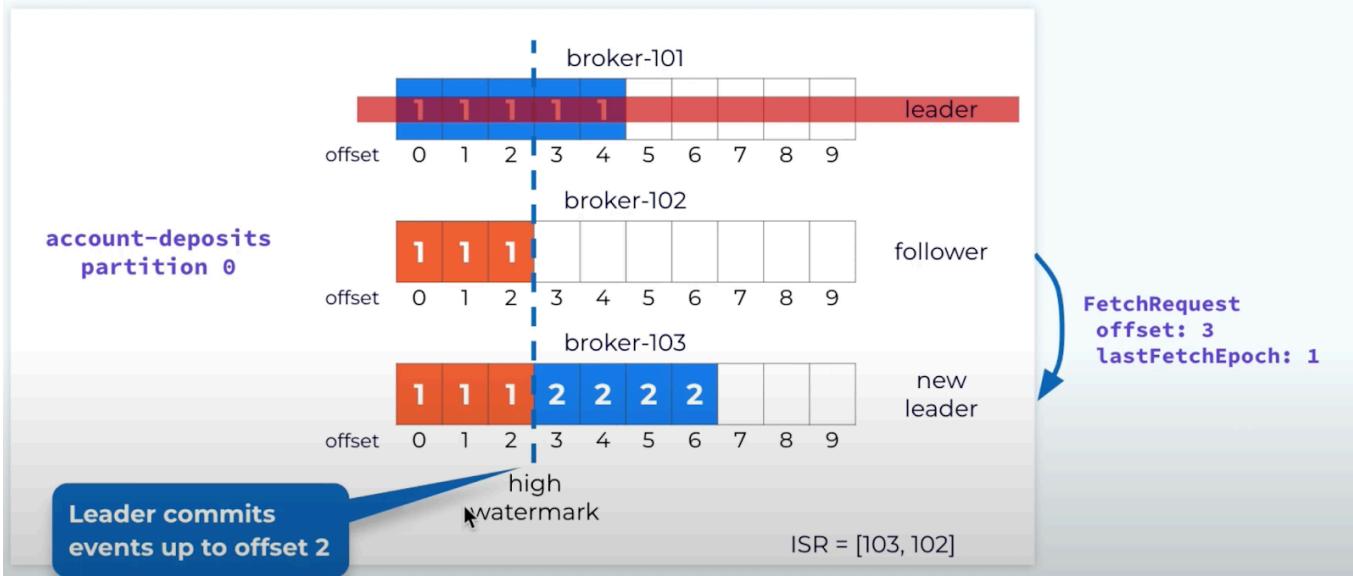
Fetch Response Informs Follower of Divergence



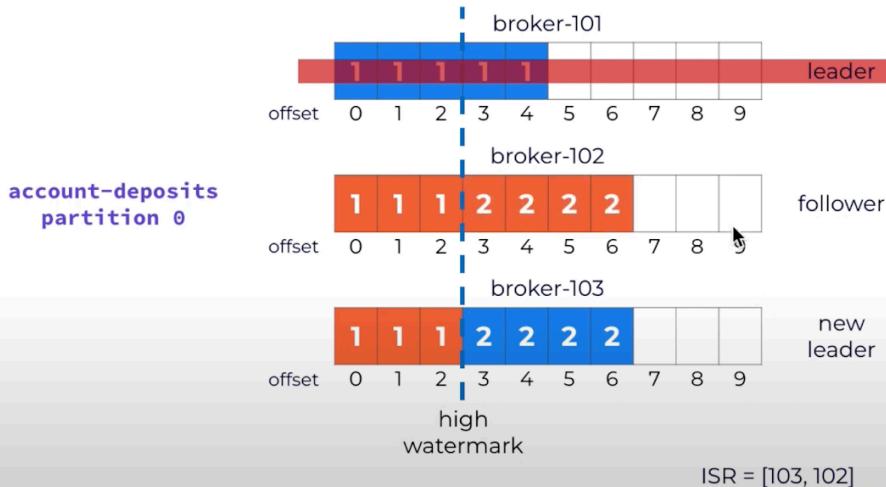
Follower Truncates Log To Match Leader Log



Subsequent Fetch with Updated Offset & Epoch



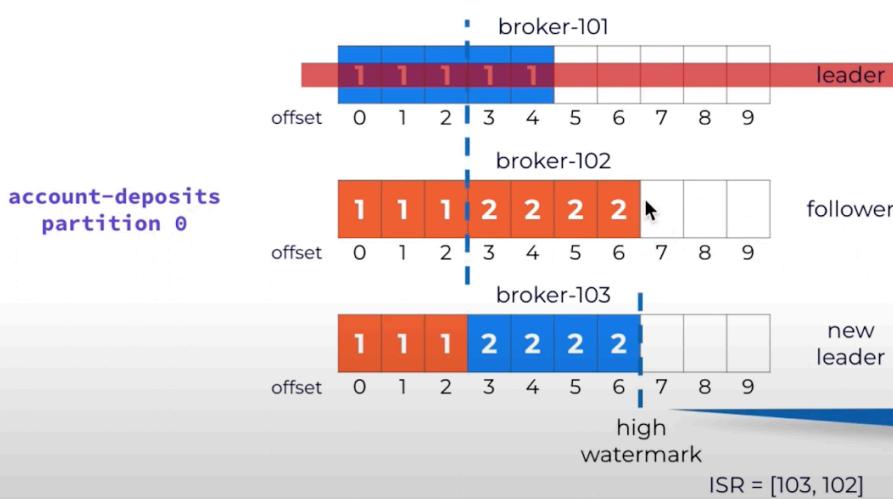
Follower 102 Reconciled



Leader sends requested offsets

FetchResponse
HighwaterMarkOffset: 3
records: [
(offset: 3, epoch: 2),
(offset: 4, epoch: 2),
(offset: 5, epoch: 2),
(offset: 6, epoch: 2)
]

Follower 102 Acknowledges New Records

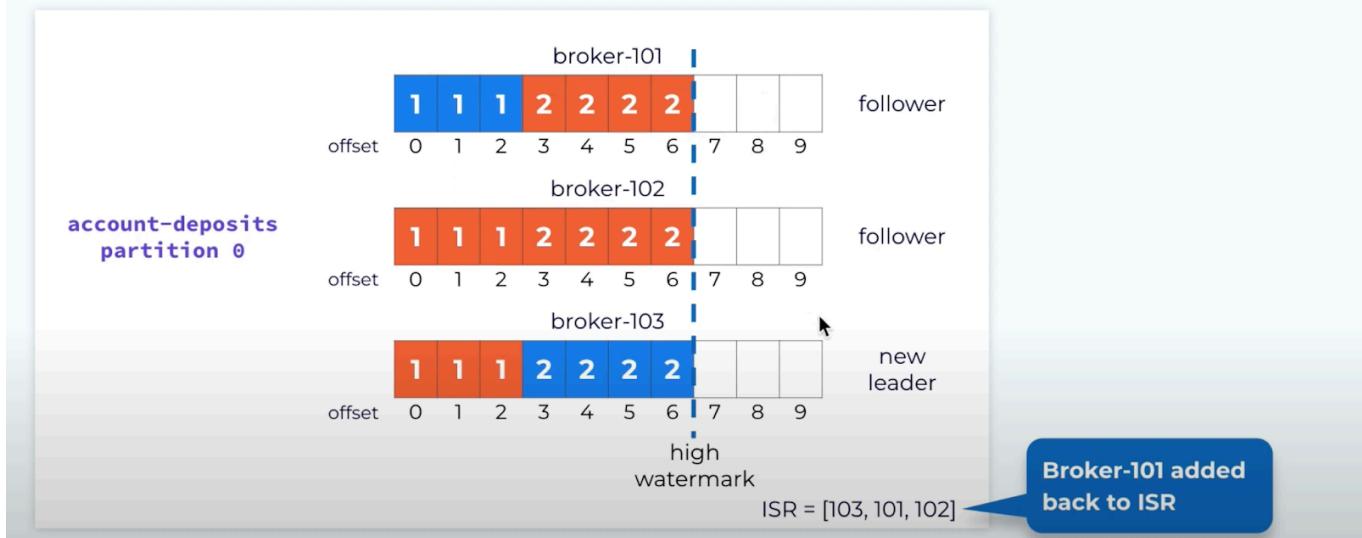


New request with latest offset and epoch values

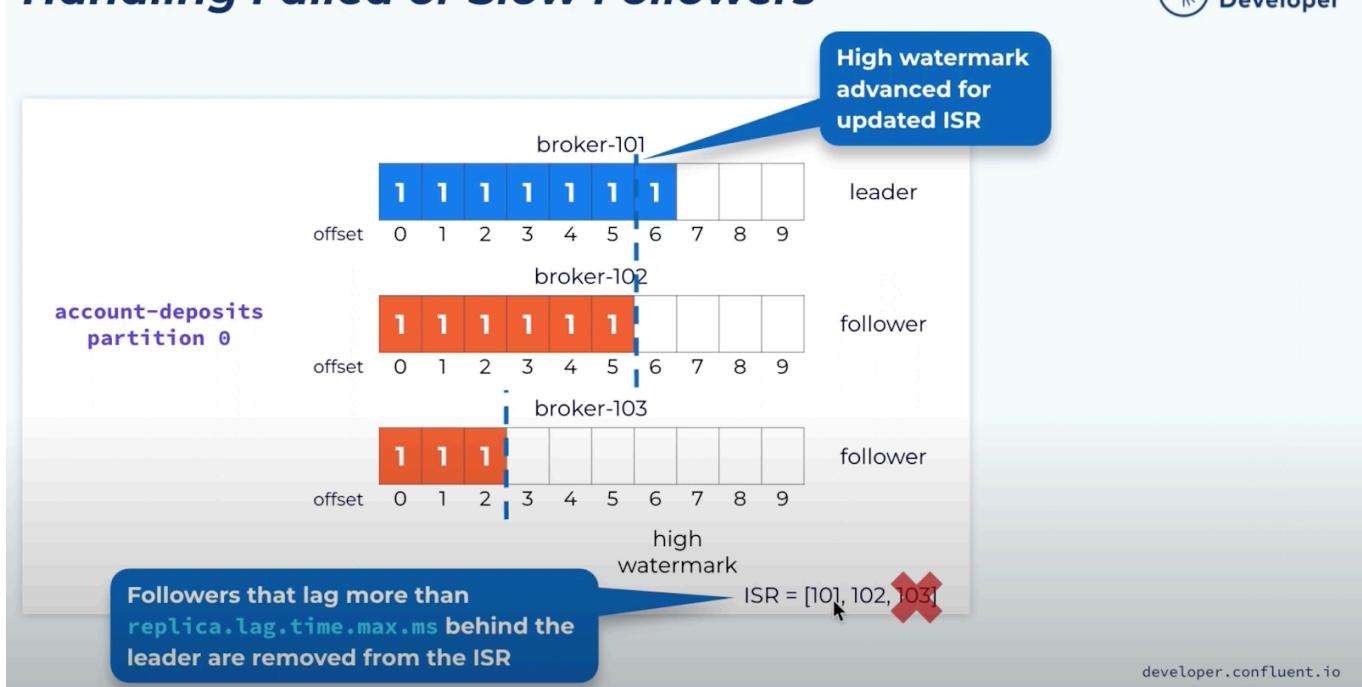
FetchRequest
offset: 7
lastFetchEpoch: 2

Leader advances its high watermark

Follower 101 Rejoins the Cluster



Handling Failed or Slow Followers



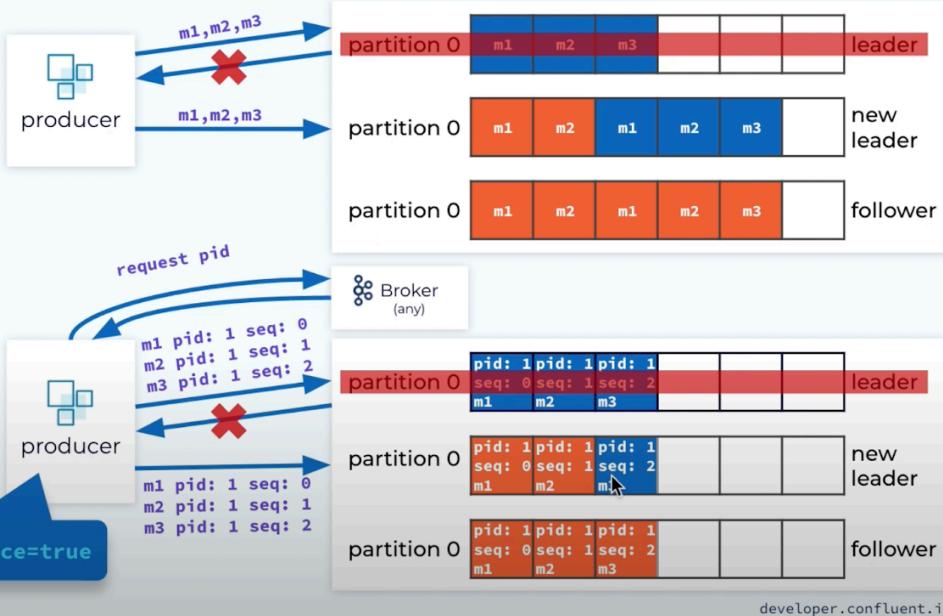
developer.confluent.io

Producer Idempotence:

Producer Idempotence



**Without idempotence,
duplicate and out of
order records can occur**

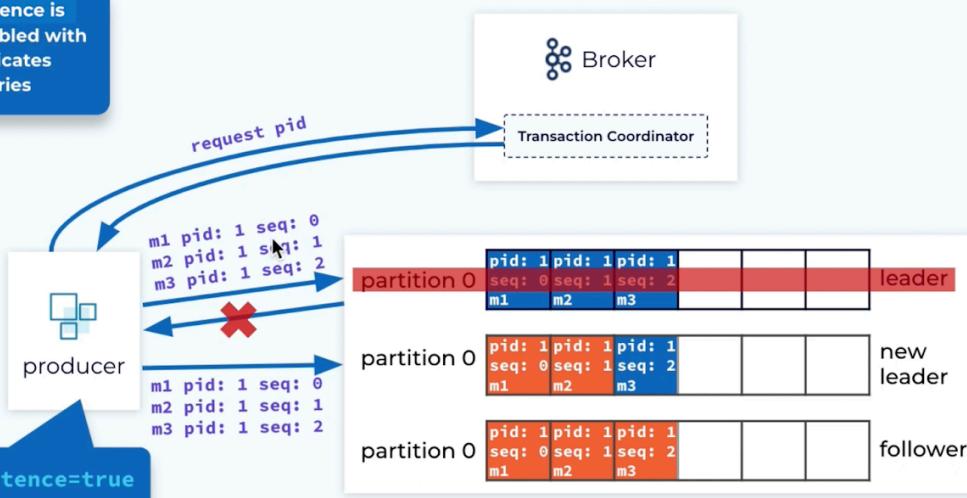


developer.confluent.io

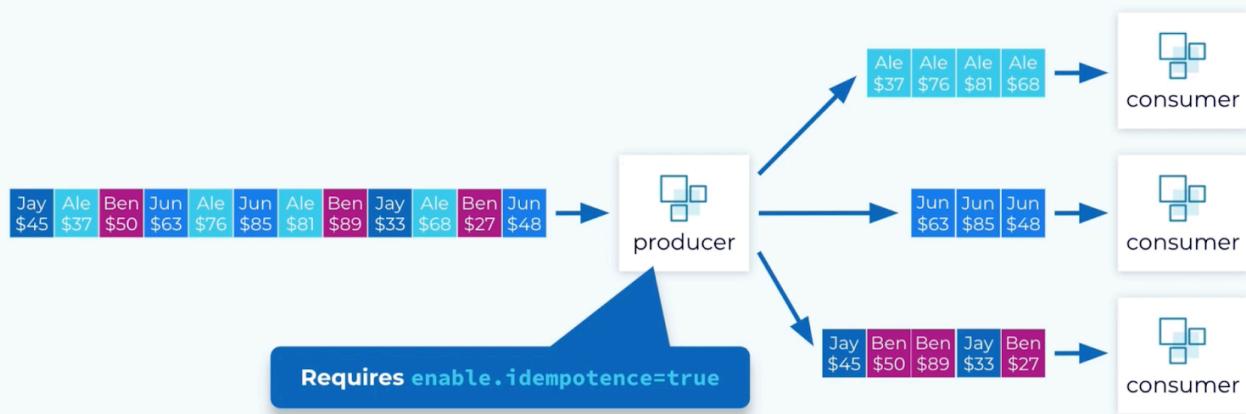
Transactions: Producer Idempotence



Producer idempotence is automatically enabled with EOS to avoid duplicates from producer retries

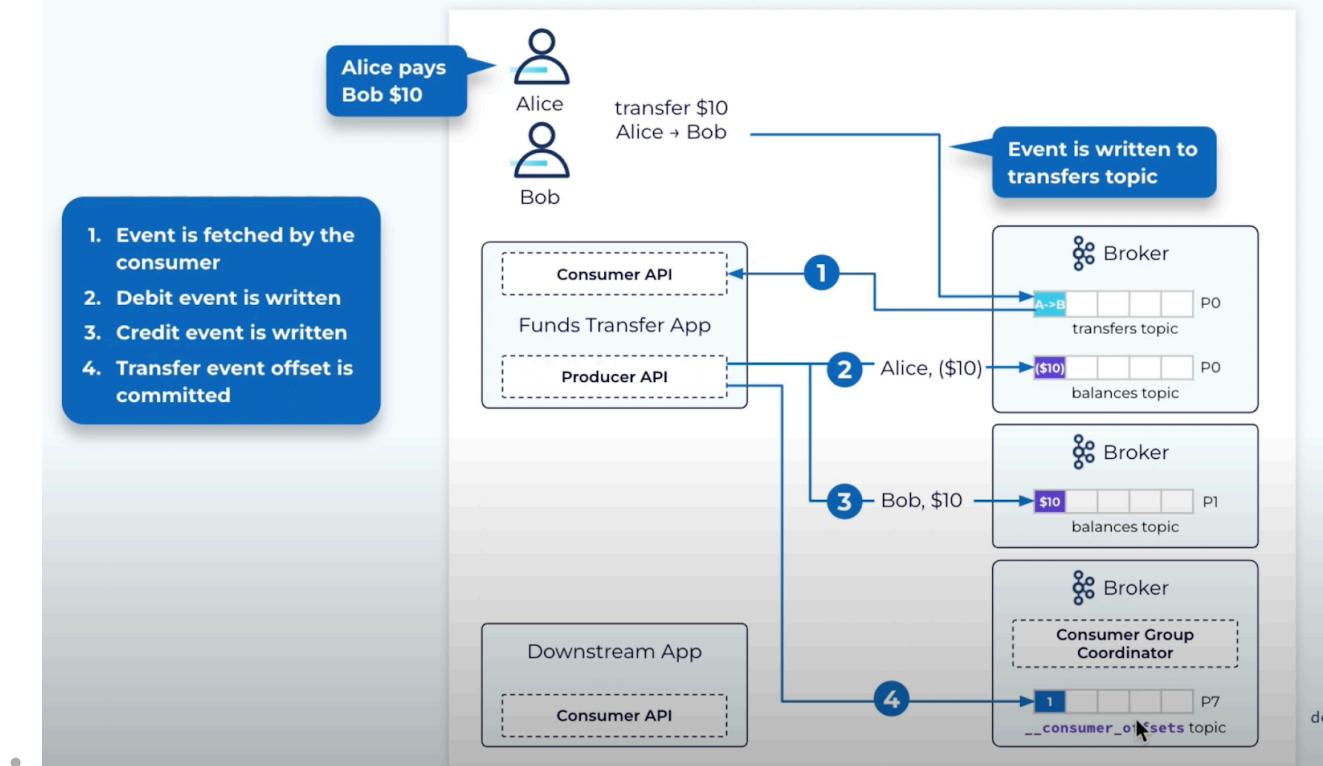


End-to-End Ordering Guarantee



Transactions:

Why Are Transactions Needed?



Kafka Transactions Deliver Exactly Once

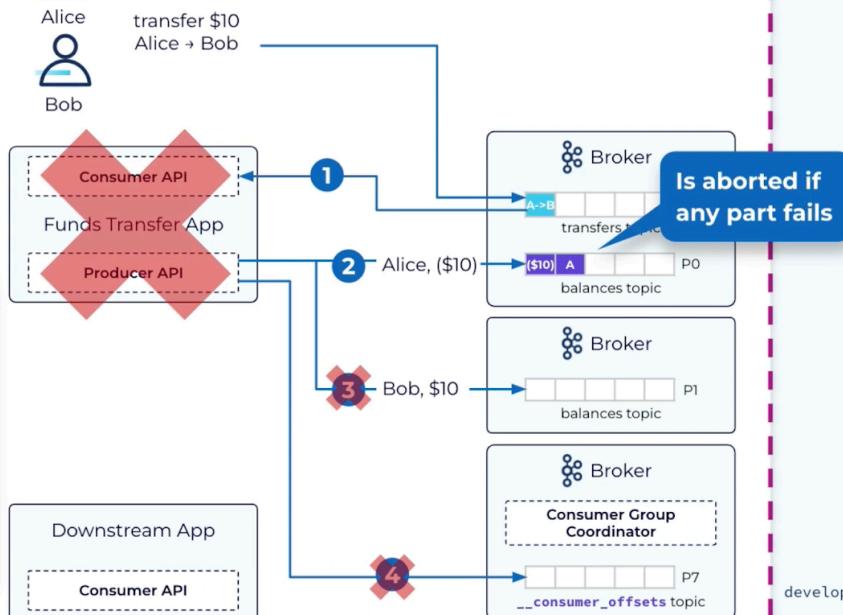


Transaction is only committed if all parts succeed

Using transactions with Kafka Streams is quite simple:

- 1) Set `processing.guarantee` to `exactly_once_v2` in `StreamsConfig`
- 2) Set `isolation.level` to `read_committed` in the Consumer configuration

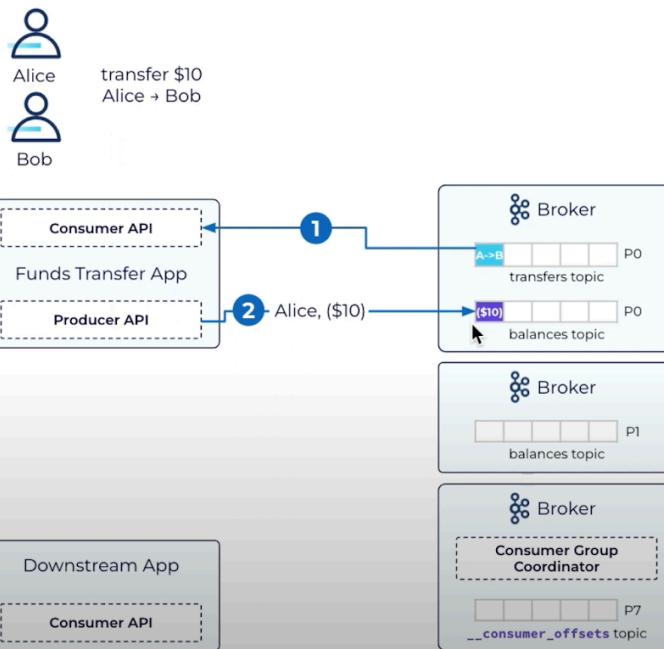
Atomic Transaction



System Failure Without Transactions



1. Event fetched by consumer
2. Alice's account debited

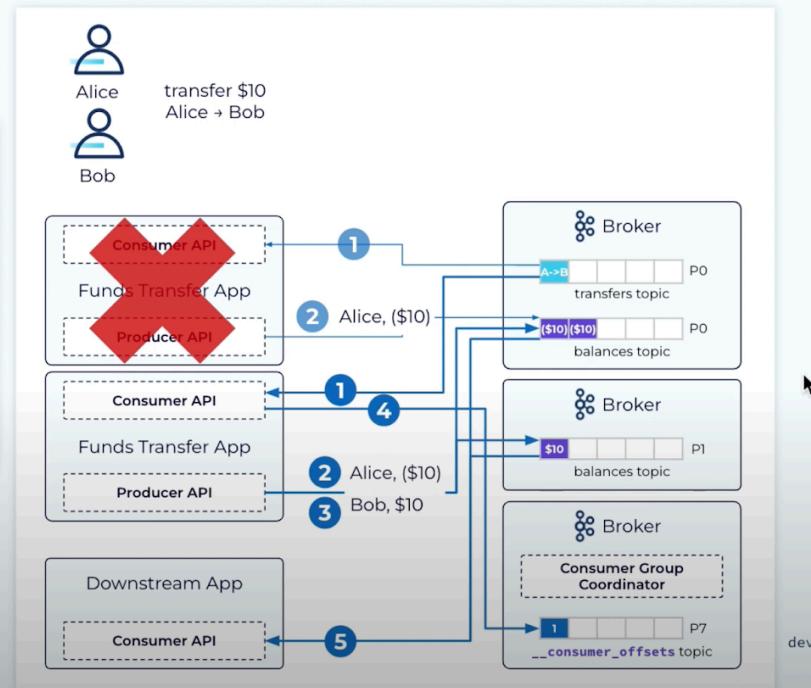


System Failure Without Transactions

1. Event fetched by consumer
2. Alice's account debited

Application instance fails without committing offset and new application instance starts

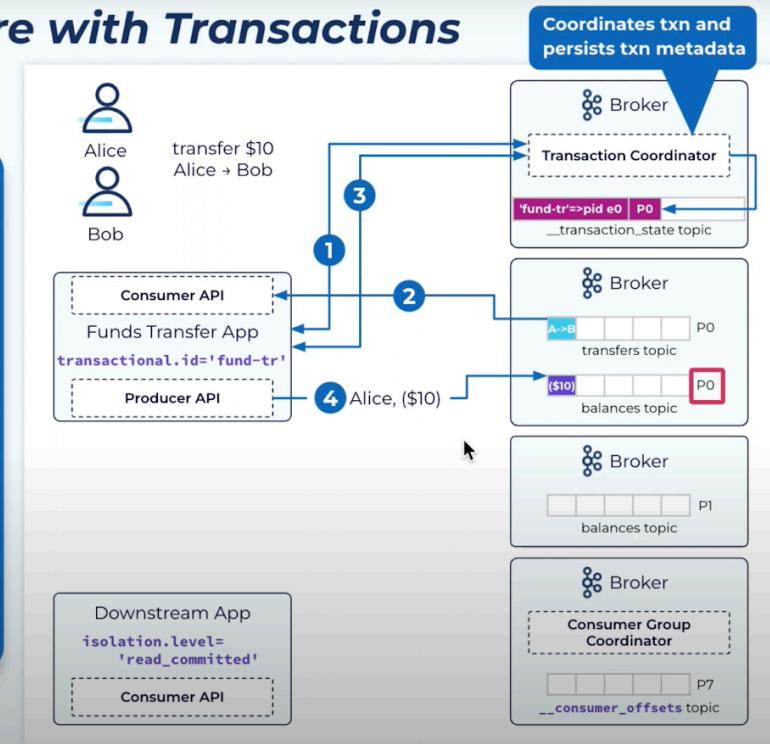
1. Event fetched by consumer
2. Alice's account is debited a second time
3. Bob's account is credited
4. Consumer offset committed
5. Two debit events processed by downstream consumer



developer.confluent.io

System Failure with Transactions

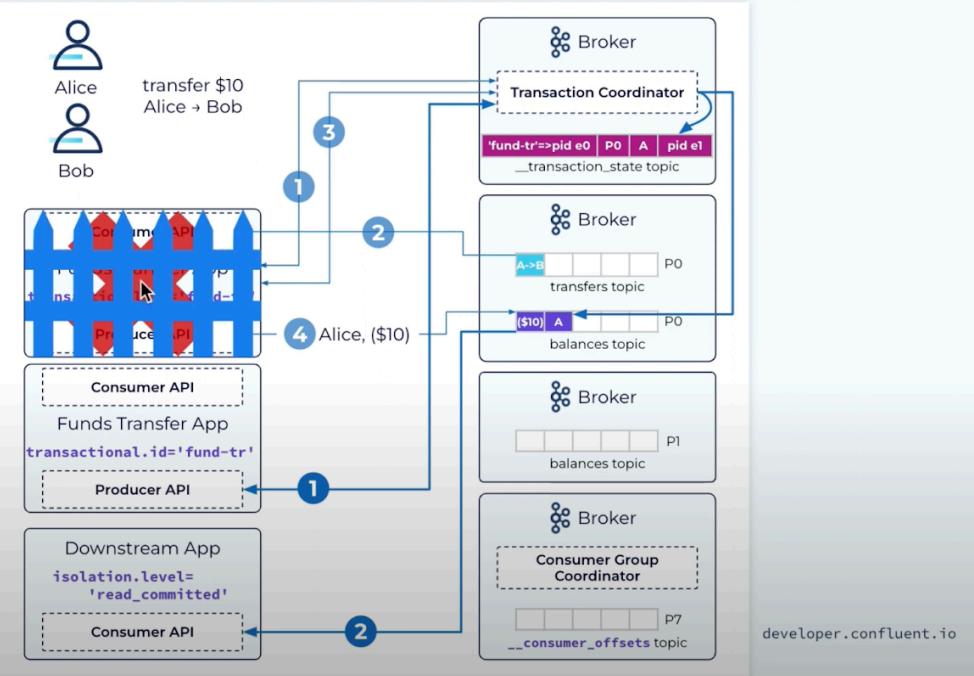
1. Requests txn ID, is returned PID and txn epoch
2. Event fetched by consumer
3. Notifies coordinator of partition being written to
4. Alice's account debited



developer.confluent.io

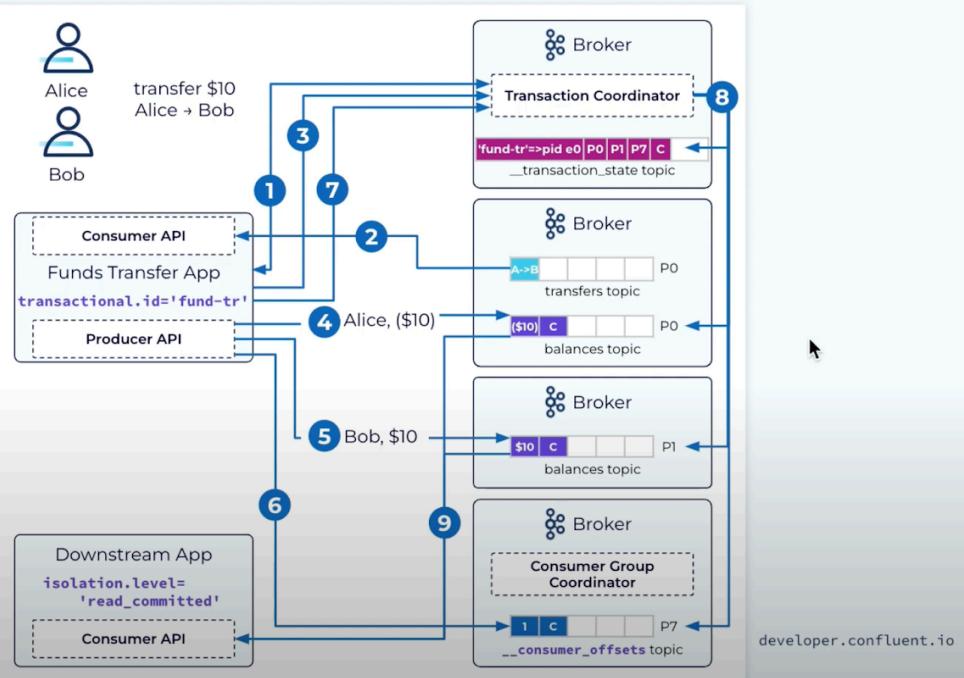
System Failure with Transactions

1. Requests txn ID, is returned PID and txn epoch
 2. Event fetched by consumer
 3. Notifies coordinator of partition being written to
 4. Alice's account debited
- Application instance fails without committing offset and new application instance starts**
1. New instance requests txn ID
 - a. Coordinator fences previous instance by aborting pending txn and bumping up epoch
 2. Downstream consumer with `read_committed` discards aborted events



System with Successful Committed Transaction

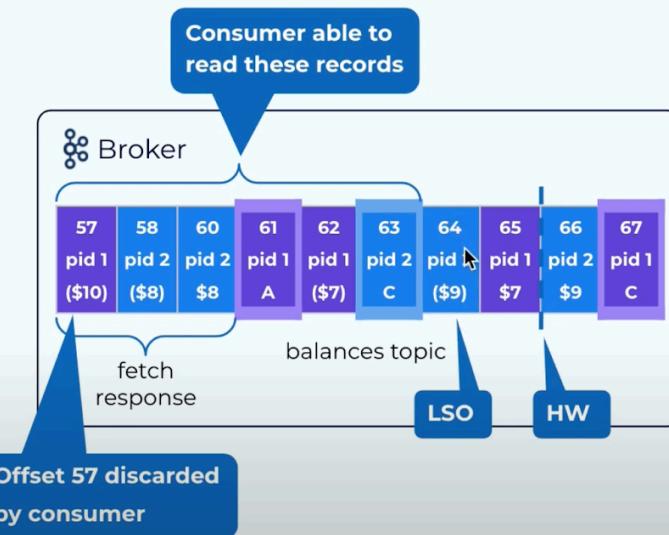
1. Requests txn ID and assigned PID and epoch
2. Event fetched by consumer
3. Notifies coordinator of partition being written to
4. Alice's account debited
5. Bob's account credited
6. Consumer offset committed
7. Notify coordinator that transaction is complete
8. Coordinator writes commit markers to p0, p1, p7
9. Downstream consumer with `read_committed` processes committed events



Consuming Transactions with `read_committed`



- Leader maintains last stable offset (LSO), the smallest offset of any open transaction
- Fetch response includes
 - only records up to LSO
 - metadata for skipping aborted records

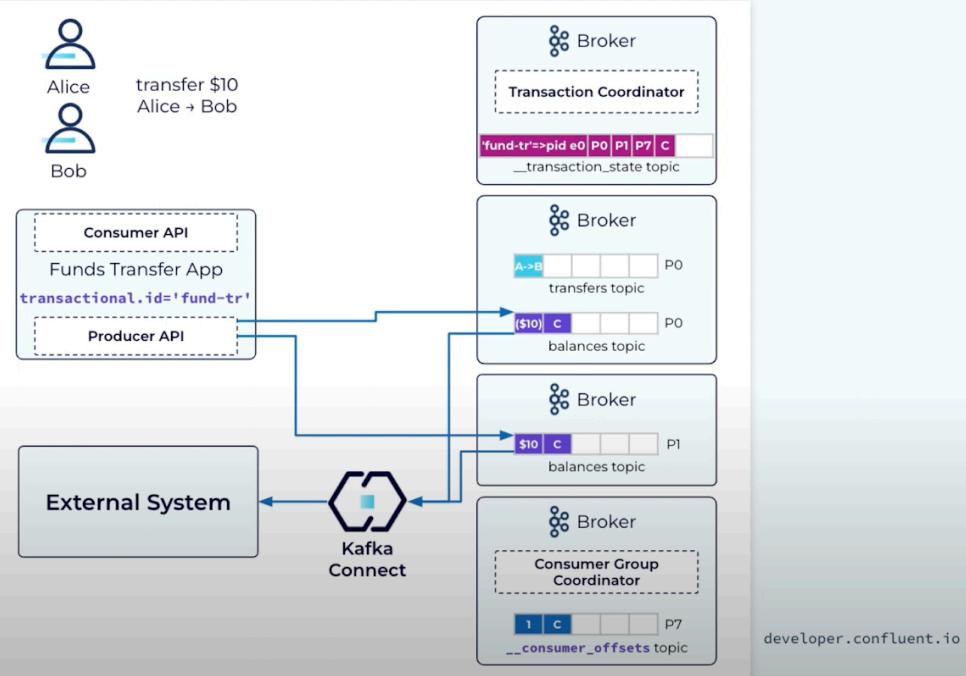


Interacting with External Systems



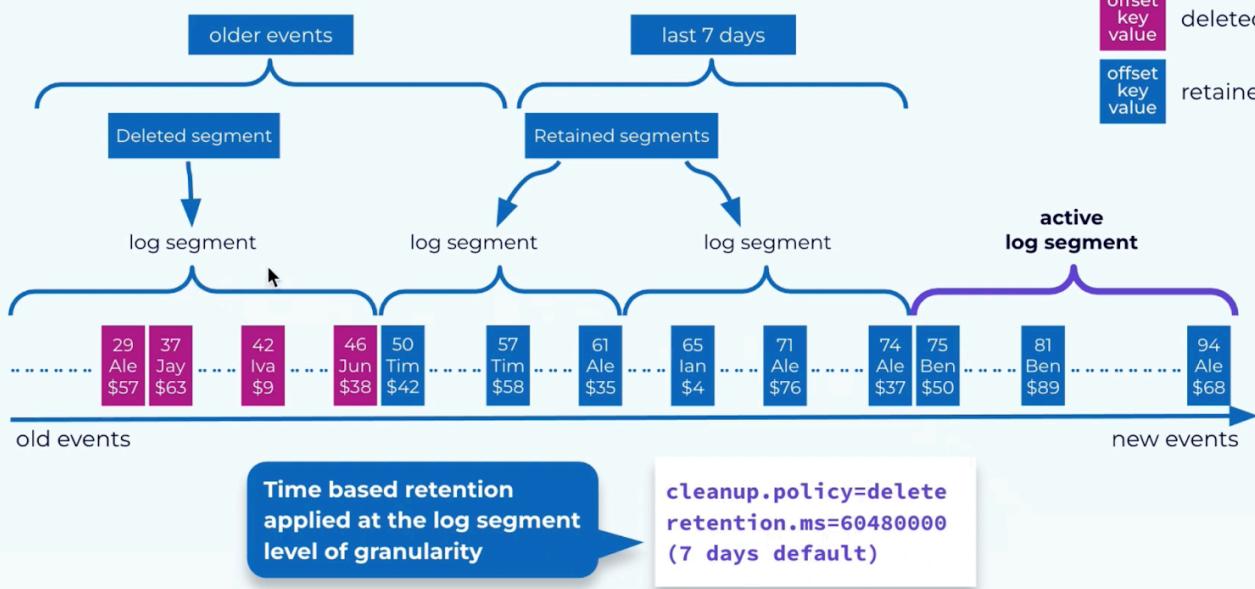
Atomic writes to Kafka and external systems are not supported

- Instead, write the transactional output to a Kafka topic first
- Rely on idempotence to propagate the data from the output topic to the external system



Retention:

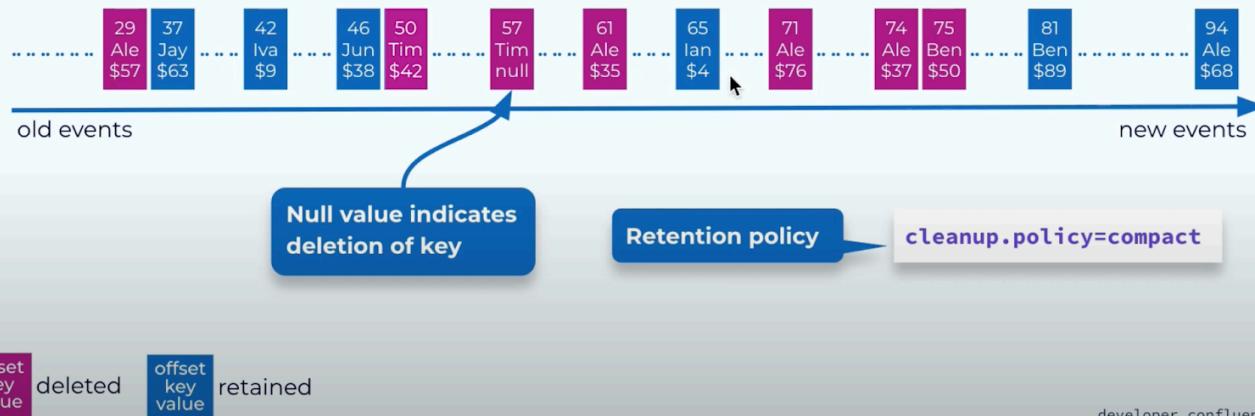
Time-Based Retention



Topic Compaction: key-based retention



- Most recent occurrence per key is kept
- All other occurrences marked for deletion over time



developer.confluent.io

Usage and Benefits of Topic Compaction

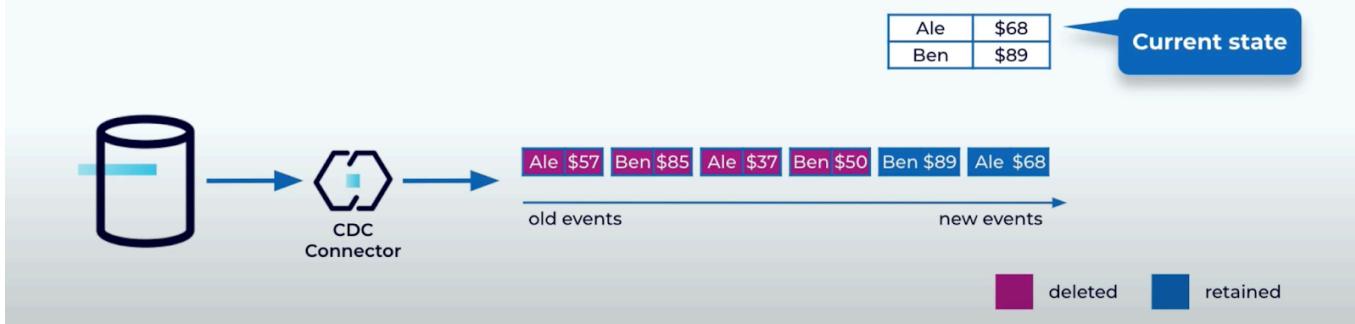


Usage:

- Streaming updatable data sets (e.g. profiles, catalogue)
- KSQL state and tables

Benefits:

- Keeping latest info with bounded storage
- Allows incremental consumption and bootstrap to be done the same way



Compaction Process - Segments to Clean



of dirty segments cleaned limited by log cleaner thread memory

Segments selected for cleaning

Active segment is never compacted



Compaction Process - Build Dirty Segment Map

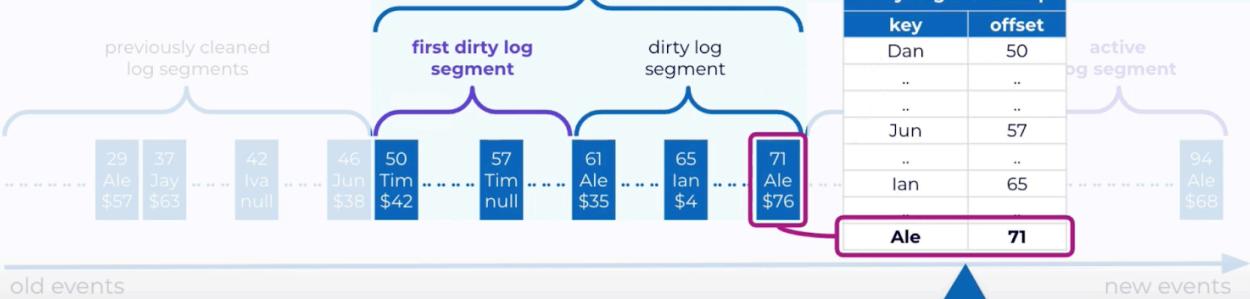


Segments selected for cleaning

Dirty Segments Map

key	offset
Dan	50
..	..
..	..
Jun	57
..	..
Ian	65
Ale	71

active log segment

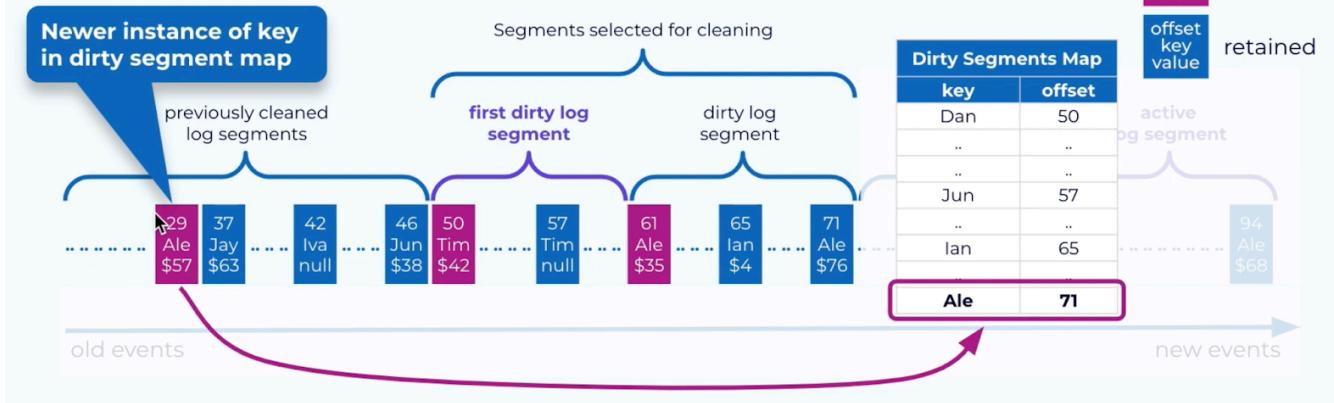


In-memory map of latest offset for each key in selected dirty segments

Compaction Process - Deleting Events



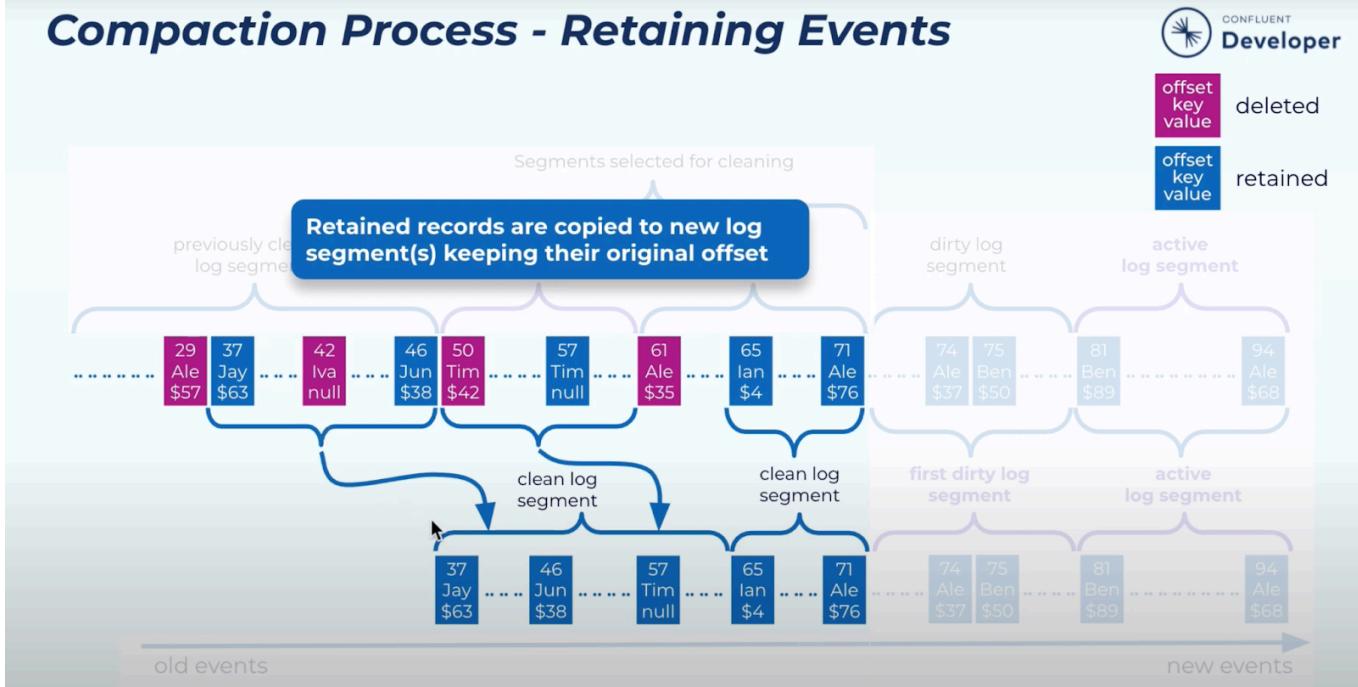
offset key value
deleted
offset key value
retained



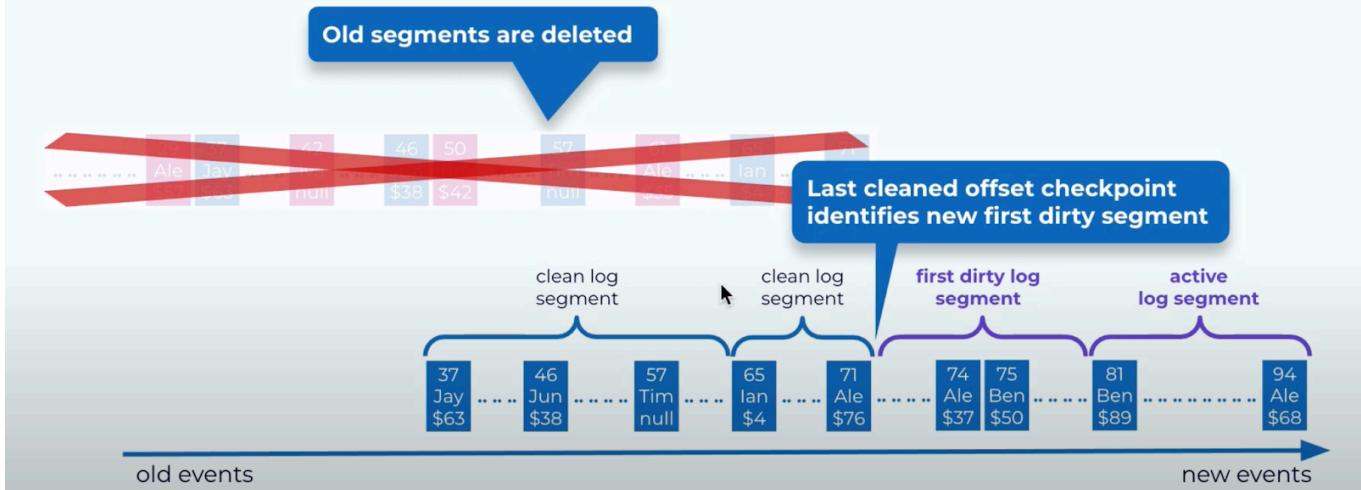
Compaction Process - Retaining Events



offset key value
deleted
offset key value
retained



Compaction Process - Replace Old Segments

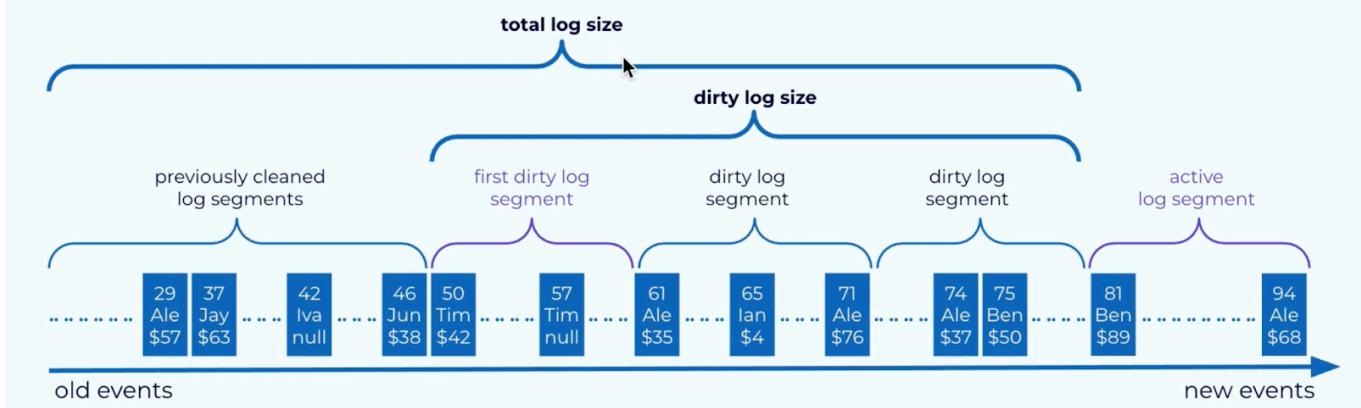


When Compaction is Triggered



A topic partition is compacted if one of the following is true:

1. `dirty / total > min.cleanable.dirty.ratio` and
`message timestamp >= current time - min.compaction.lag.ms`
2. `message timestamp < current time - max.compaction.lag.ms`



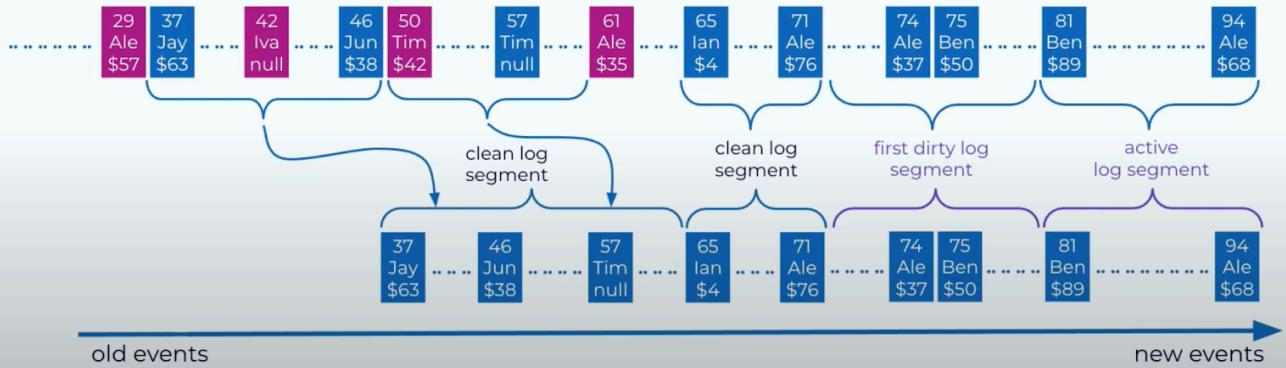
Topic Compaction Guarantees



A consumer:

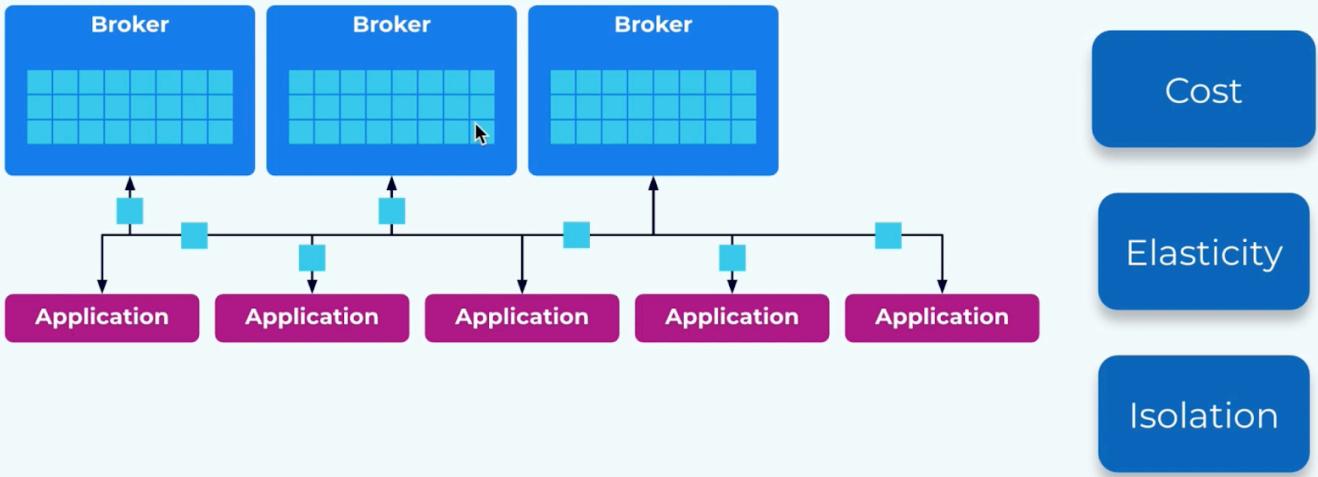
- 1) is not guaranteed to see every record
- 2) is guaranteed to see the latest record for a key

offset key value	deleted
offset key value	retained



Kafka tiered storage:

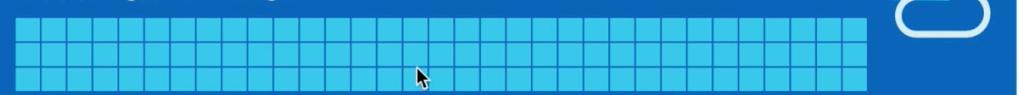
Current Kafka Storage and Its Shortcomings



Tiered Storage

Remote Object Storage

AWS S3, Google Cloud Storage



Cost
Efficiency

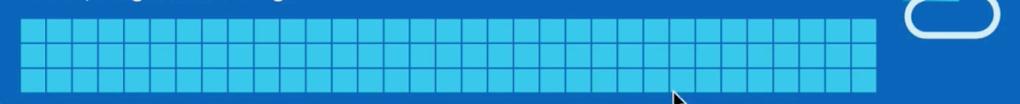
Elasticity

Isolation

Tiered Storage

Remote Object Storage

AWS S3, Google Cloud Storage



Cost
Efficiency

True
Elasticity

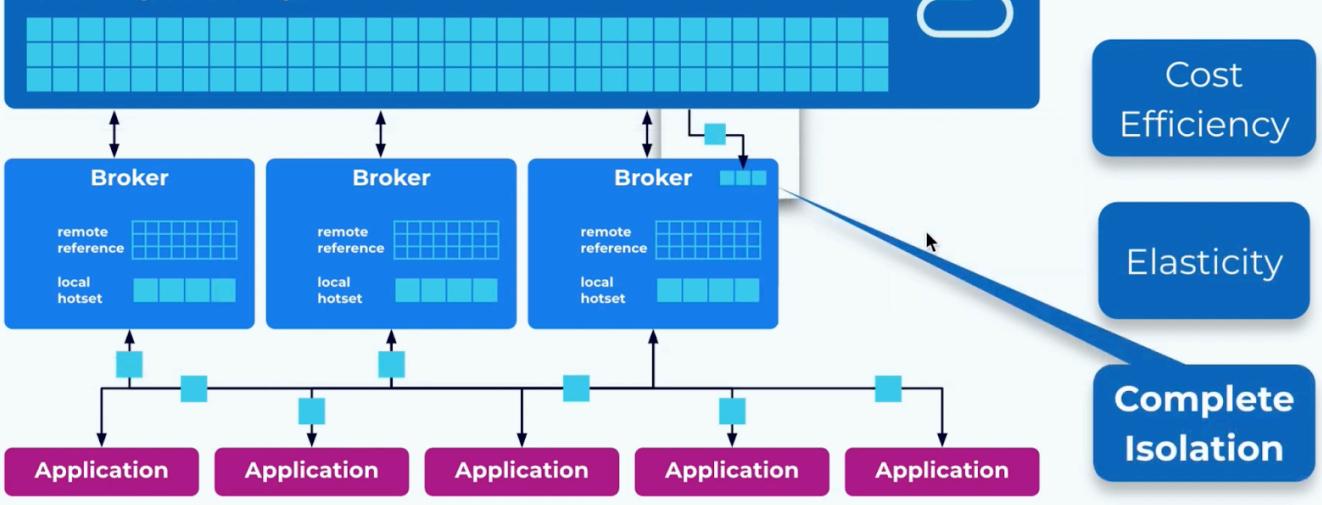
Isolation

Tiered Storage



Remote Object Storage

AWS S3, Google Cloud Storage



Fetching Tiered Data



If data is in the hotset, it is returned using standard Kafka methods

